

Ayushkant Behera

Phone: +91 9124706434 | Email: ayushkantworks@gmail.com

[LinkedIn](#) | [GitHub](#) | [Portfolio](#) | [Medium](#)

PROFESSIONAL SUMMARY

Full-stack developer who builds real, working products — AI agents and computer-vision automation among them. I also freelance solo, handling everything from architecture to shipping.

PROFESSIONAL EXPERIENCE

Founder & Developer — rizzUP [Live Site](#) [GitHub](#)

- Built and shipped rizzUP, an AI dating reply generator, owning product design, LLM pipeline engineering, and payments infrastructure end-to-end.
- Used a vision model (Qwen3-VL) to read and classify conversations (icebreaker vs. thread reply, tone, context) so generated replies are grounded in the actual back-and-forth rather than the raw text alone.
- Built a multi-stage generation and retry pipeline with programmatic scoring to catch flat lines, wrong language, or generic responses before they reach the user.
- Shipped a live subscription product with Dodo Payments checkout, signed webhooks, and a server-side paywall — grew to 20 paying subscribers at \$15/month through TikTok marketing.

Freelance Full Stack Developer

- Built two paid iOS game-automation bots — a Triumph.gg blackjack bot ([GitHub](#)) using computer vision for card/state detection, and a word-link solver ([GitHub](#)) using OpenCV template matching and DFS — \$250 each.

Full Stack Developer Intern — LevelUp Labs [Live Site](#)

- Contributed to the development of a structured learning platform that organizes YouTube educational content based on user learning goals.
- Implemented authentication and user management using Clerk.
- Integrated Supabase as backend infrastructure and database solution.
- Improved UI/UX and implemented progress tracking features for learners.

PROJECTS

rizzUP – Context-Aware AI Reply Assistant

Next.js · Hugging Face Inference API · Supabase · Dodo Payments

[GitHub](#) | [Live Site](#)

- Built a web app that turns "what do I even say?" into one context-aware reply — users paste a message or screenshot of a real conversation and get back a genuinely relevant response instead of generic flirty suggestions.
- Used a vision model (Qwen3-VL) to read and classify the conversation (icebreaker vs. thread reply, tone, context) before generating a reply, so output is grounded in the actual back-and-forth rather than the raw text alone.
- Implemented a multi-stage generation and retry pipeline with programmatic scoring to catch flat lines, wrong language, or generic responses before they reach the user.
- Shipped a live subscription product with Dodo Payments checkout, signed webhooks, and a server-side paywall — grew to 20 paying subscribers at \$15/month through TikTok marketing.

SelfHeal

Next.js · TypeScript · Groq · OpenAPI · Tailwind CSS

[GitHub](#) | [Live Site](#)

- Self-correcting API agent — turns a plain-English goal into a real HTTP request against any public OpenAPI/Swagger-described API, executes it, and self-heals on failure by feeding the error back to an LLM and retrying (up to 3 attempts), streaming every attempt live to a terminal-style UI.
- Designed a keyword-based endpoint pre-filtering step to keep large specs (e.g. GitHub's 1,220-endpoint API) within LLM rate limits — reduced a 252KB spec summary to ~3KB via position-weighted path/keyword matching before the LLM ever sees it.
- Enforced a strict execution boundary: the LLM only ever outputs structured JSON (method/url/headers/body); a fixed server-side fetch() is the sole thing that touches the network — no eval, no arbitrary code execution.

Across – Browser Memory Chrome Extension

Manifest V3 · RAG · pgvector

[GitHub](#) | [Watch Demo Video](#)

- Built a Chrome extension that continuously indexes open browser tabs and provides a natural-language chat interface grounded in the user's browsing history using RAG (Retrieval-Augmented Generation).
- Automatically tracks tab lifecycle events and extracts readable page content via Mozilla Readability, splitting pages into heading-aware semantic chunks with overlap.
- Stores vector embeddings server-side in PostgreSQL with pgvector, supporting multiple LLM providers (Groq Llama 3, OpenAI GPT-4o-mini, Anthropic Claude) and embedding providers (Jina AI, OpenAI, Hugging Face).
- Delivers source-cited answers grounded in content the user has actually viewed; works completely free using Jina AI embeddings and Groq Llama 3 with no credit card required.

EDUCATION

Manipal University Jaipur — Bachelor of Computer Applications (BCA)

2028

- Pursuing a flexible-format degree, allowing full-time availability for freelance work, internships, and full-time roles alongside coursework.

SKILLS

Languages: TypeScript, JavaScript, Python, C# (Unity)

AI / LLM Engineering: RAG pipelines, LLM integration (OpenAI, Anthropic, Groq), embeddings

Frameworks & Libraries: React, Next.js, TanStack Start/Query, Node.js, Express.js, Tailwind CSS

Databases: PostgreSQL (pgvector), MongoDB, Supabase

Automation & Computer Vision: OpenCV, WebDriverAgent (WDA), screen-capture pipelines, template matching

Game Development: Unity 6.3 LTS, gameplay/event-system architecture

Security & Reverse Engineering: API/auth flow analysis, traffic interception, MITM proxy debugging

Tools: Git, Linux, Clerk, Vercel